

Attached here is a link to zip containing a data file containing all pitches from the 2011 MLB season and the associated metadata describing the data.

Please review the dataset and start outlining the way you would go about the model-building process with the goal of predicting the probability that the next thrown pitch will be a fastball, slider, change-up, etc.

We would like to see you assess the data, build a predictive model, and then evaluate that model and the success of its predictions, with the understanding that delivering predictions with any degree of accuracy is unlikely in this short time span. Please provide any associated data analysis (plots, graphs, etc.), feature engineering, and code assembled in the form of a python notebook. Please include markup text to explain your analysis, graphs, etc. If you include a notebook, please also add an HTML version of that notebook, to facilitate review.

We would like you to spend no more than 4-6 hours on this project so please provide some details around future steps you would take if you had more time to finalize this project. This is meant to be an open-ended/creative project, so focus your time where you see best. Let me know if you have any questions along the way. We aim to have a turnaround time to return the project in a week. We also understand that this isn't always possible, so if you need additional time, please let me know so I can keep the Hiring Team in the loop.